





What this shows: each arrow starts at the model's unsteered Yes/No bias and ends after adding the learned behavioral-factuality direction before the answer. Arrows can point differently in empty and conversation contexts, showing that the same intervention may have context-dependent causal effects.