



Misinformation induction

Unrelated-topic conversation

(a) Internal representation of the original claim



(b) Accuracy of direct prompts about the original claim



What this shows: the original fact is measured at baseline, while fabricated information is reinforced, and at the configured checkpoints after the conversation switches to an unrelated topic. The upper panel uses the same internal factuality probe for comparison; the lower panel uses actual generated Yes/No answers.